

CSE 527: Introduction to Computer Vision

Week 5 - Class 1: Matching, Stitching, Registration

September 26th, 2017

••••○ Verizon LTE

3:02 PM

42%

Albums

sheepdog or mop

Select



••••○ Verizon

4:20 PM

76%

Albums

chihuahua or muffin

Select



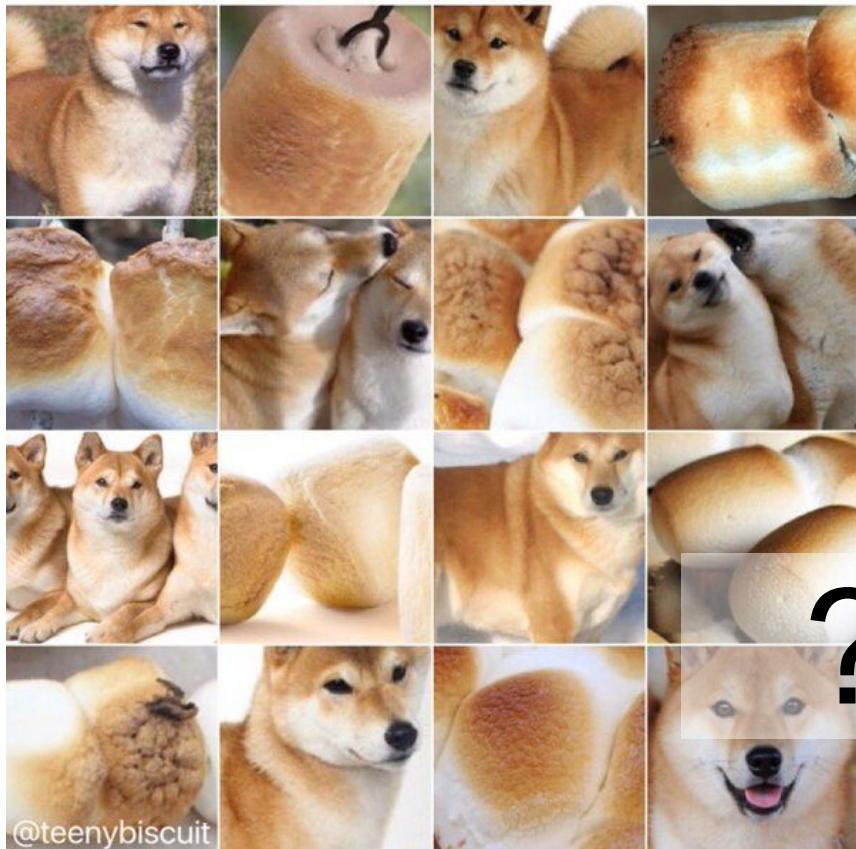
@teenybiscuit

••••○ Verizon ⌂

4:20 PM

69% ⌂

< Albums shiba or marshmallow Select



@teenybiscuit

••••○ Verizon ⌂

4:20 PM

69% ⌂

< Albums kitten or ice cream Select



@teenybiscuit

???

Recap

Gradient Histogram

A very powerful idea that is the basis for many image & feature descriptors.

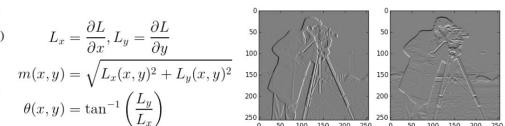
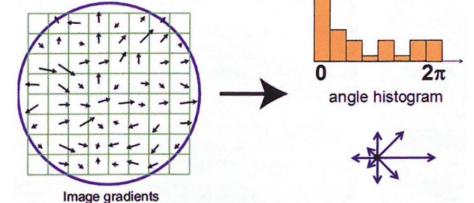
Why? Hint: invariant in more than one way...

Look at a (local) region and calculate local gradients, then bin the gradients in a histogram.

How to get **direction** and **magnitude** of edges?

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$
$$\theta(x, y) = \tan^{-1}((L(x, y+1) - L(x, y-1)) / (L(x+1, y) - L(x-1, y)))$$

Most can be pre-calculated over the entire image with a single convolution:



Feature Matching

Given features of image 1 and of image 2, how do we find the best matching of the two sets?

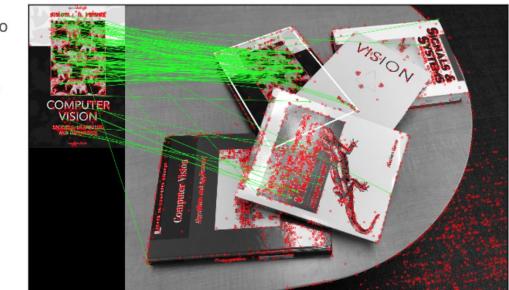
- Define a **distance metric** between a pair of feature descriptors
- Find the best **matching**.

Some metrics are determined by the descriptor (Hamming, L_2 norm).

$$M(f_L, f_R) = \|f_L - f_R\|_{L_2} < t$$

What is an **optimal matching**?

Naive: Pick the nearest neighbor (NN), set a threshold on distance.



SIFT matching and object finding

Today

Feature Matching

Image Alignment

Panoramas

HW2!

Feature Matches



Feature Matching

Given features of image 1 and of image 2, how do we find the best matching of the two sets?

- Define a **distance metric** between a pair of feature descriptors
- Find the best **matching**.

Naive: Pick the **nearest neighbor** (NN), set a threshold on distance.

$$M(f_L, f_R) = \|f_L - f_R\|_{L_2} < t$$

Apply: Ratio test and Reciprocity

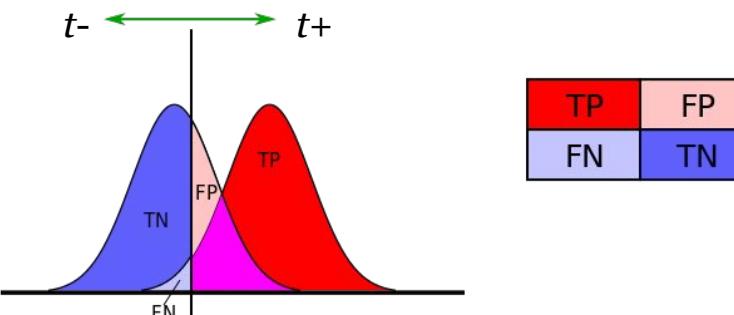


SIFT matching, ratio, recip., and object finding

Feature Matching | Measuring Performance

Receiver Operator Characteristic (ROC)
Curve

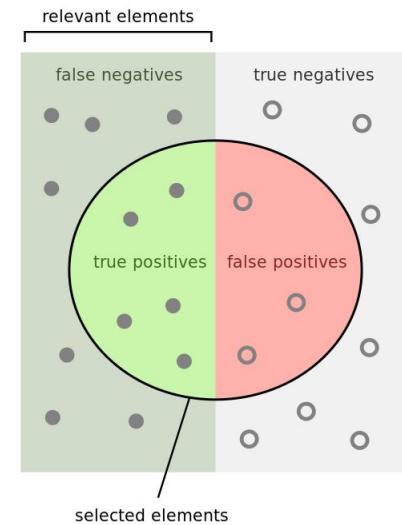
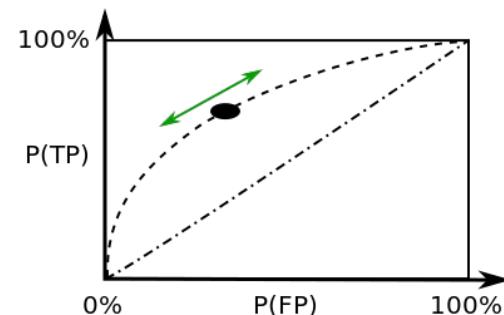
- Vary the threshold t
- Calculate precision and recall
- Take a t that has highest TP-rate, lowest FP-rate.



Precision: $TP / (TP + FP)$

"

Recall: $TP / (TP + FN)$



How many selected items are relevant?

$$\text{Precision} = \frac{\text{relevant}}{\text{selected}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{relevant}}{\text{relevant}}$$

Feature Matching | Brute Force

For every feature i in image 1

- Go over all features in image 2
- Take the one (k) that are closest to i

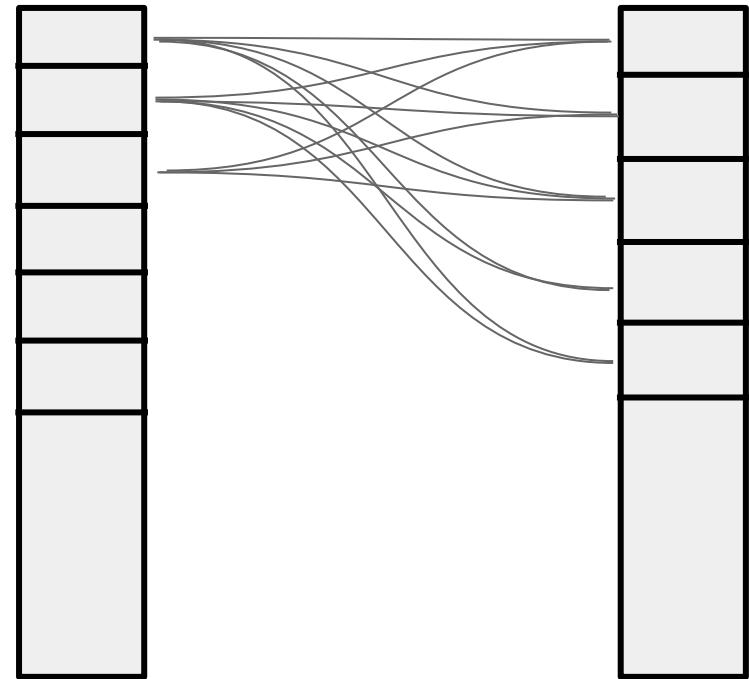
Complexity: $O(n^2)$

Good:

- Simple
- Easily parallelizable
- Exhaustive, deterministic

Bad:

- Slow



Feature Matching | k-d tree

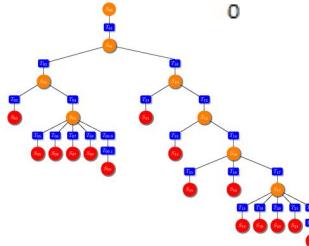
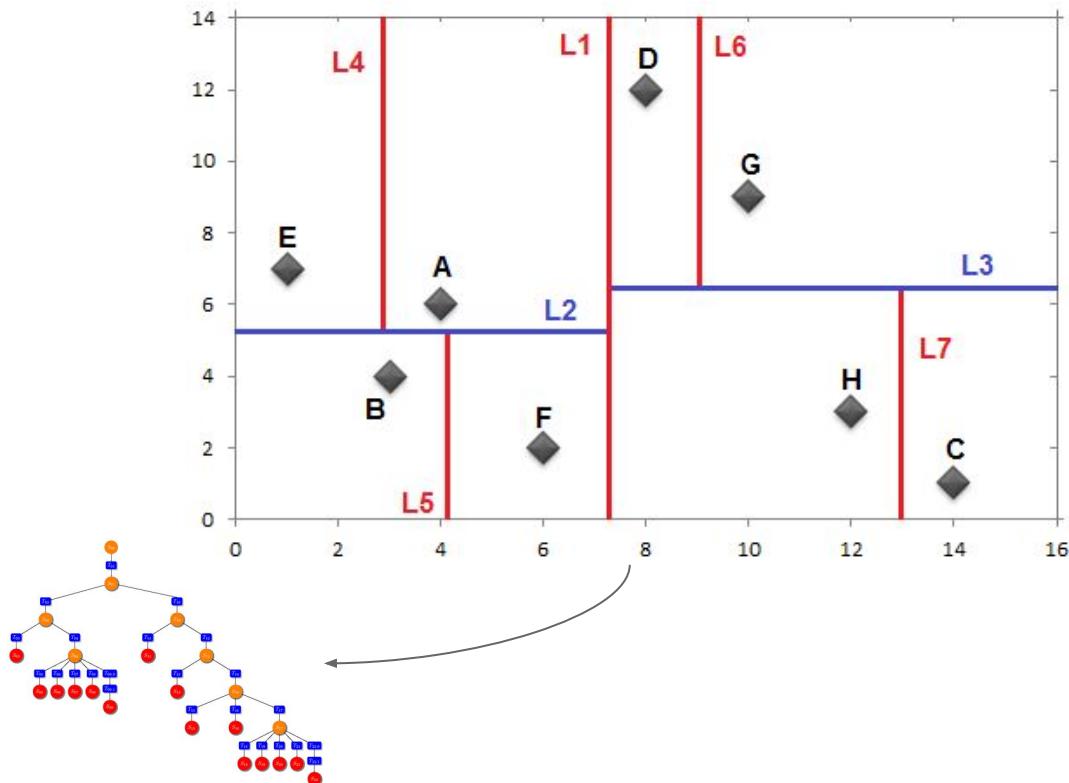
Accelerate the matching process.

Preprocess the “training” dataset:

- Hierarchically subdivide the k -dimensional space (by sorting)
- If a bin has more than m points - split to smaller spatial bins.

Build the tree: $O(k n \log n)$

Search the resulting tree: $O(\log n + m)$



Feature Matching | k-d tree

Accelerate

```
$ %timeit -n50      matches2tol = bf.knnMatch(des2,des1,k=2)  
50 loops, best of 3: 17.1 ms per loop
```

Preprocess

- Hierarchical spatial search
- If a small tree

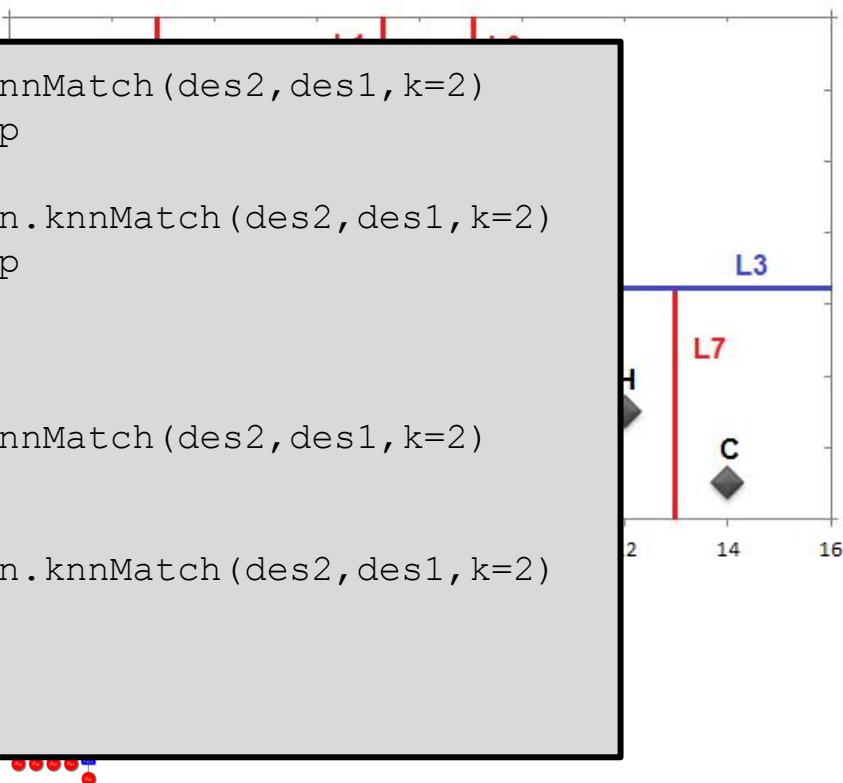
```
$ %timeit -n50      matches2tol = bf.knnMatch(des2,des1,k=2)  
50 loops, best of 3: 148 ms per loop
```

Build the tree

Search the tree

```
$ %timeit -n50      matches2tol = flann.knnMatch(des2,des1,k=2)  
50 loops, best of 3: 33 ms per loop
```

Speedup: ~ x4.4 with ~3300 features



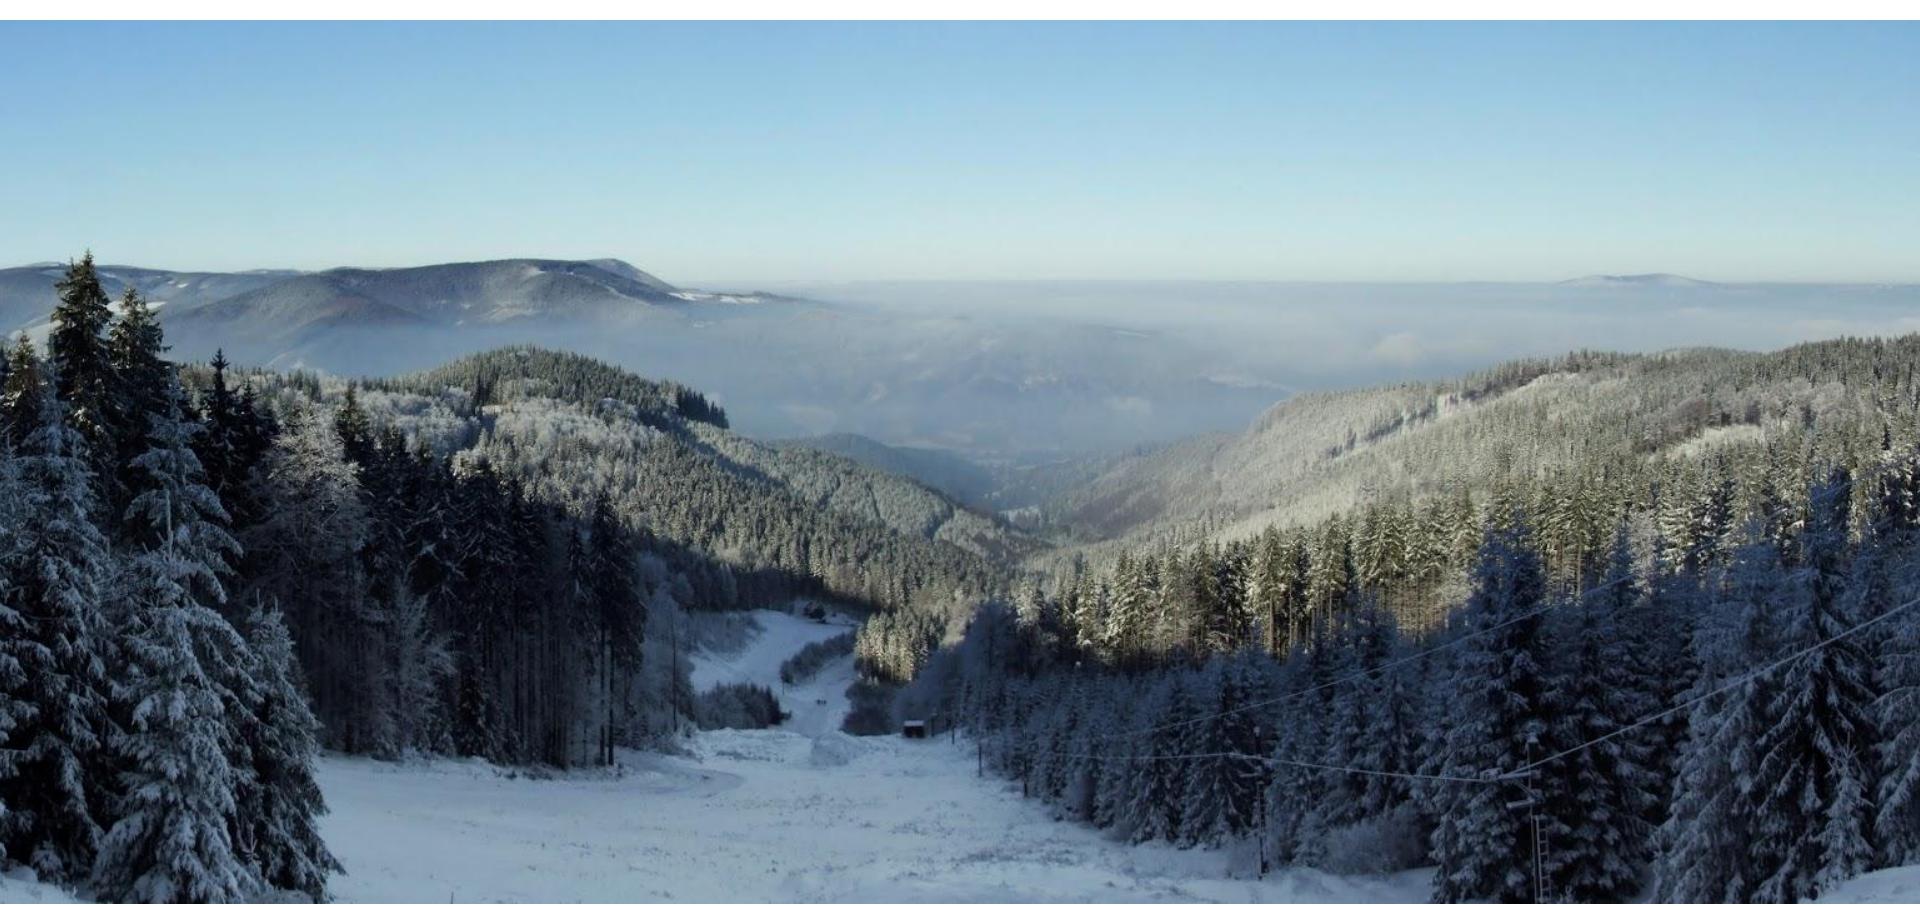
Questions?

Image Alignment





Panoramic view, right?



Panoramic view, right?



Panoramic view, right?

Image Alignment | Motivation

Why align, stitch images together?

- The human eye has $\sim 210^\circ_h \times 150^\circ_v$ FOV
- Our images/cameras have narrow FOV
- Object is just. too. big.
- We want to deliver the “big picture”



Image Alignment | Applications

Panoramas



Image Alignment | Applications

Google Street View

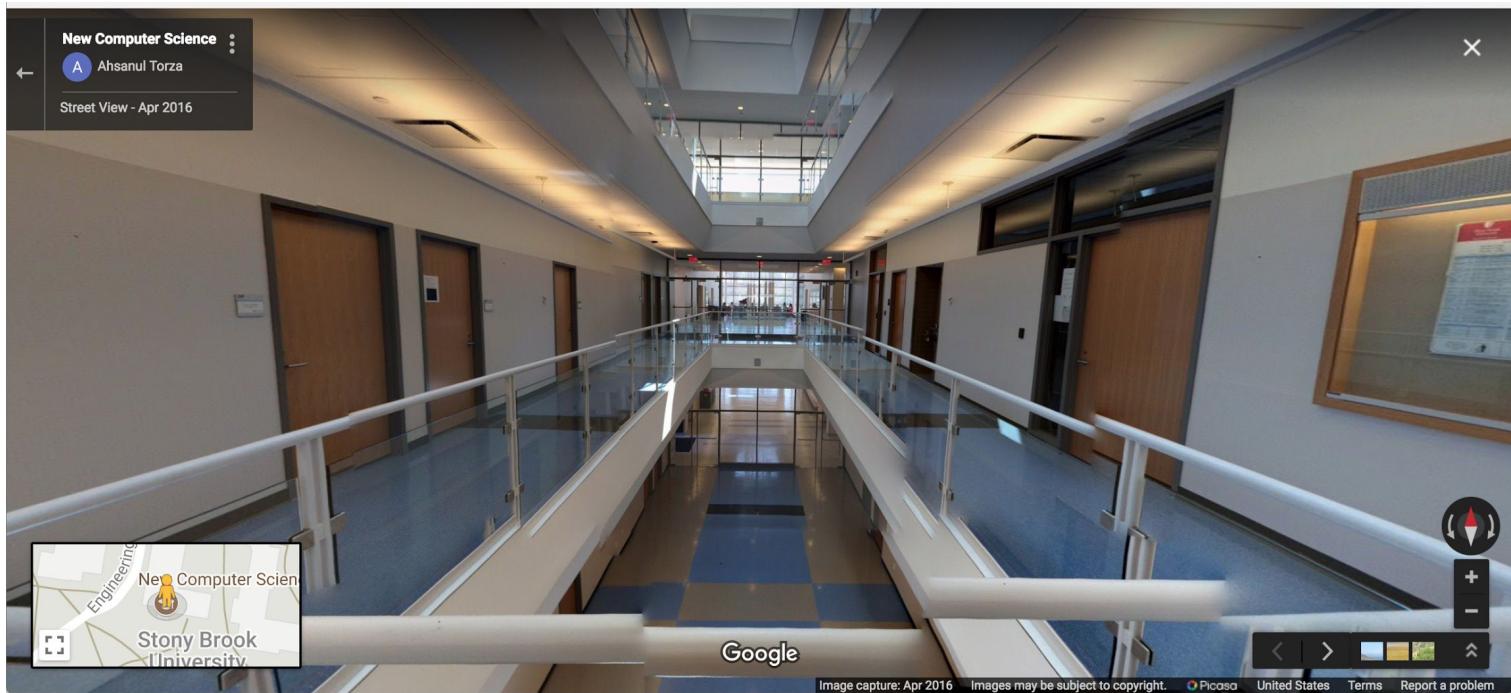


Image Alignment | Applications

360 VR

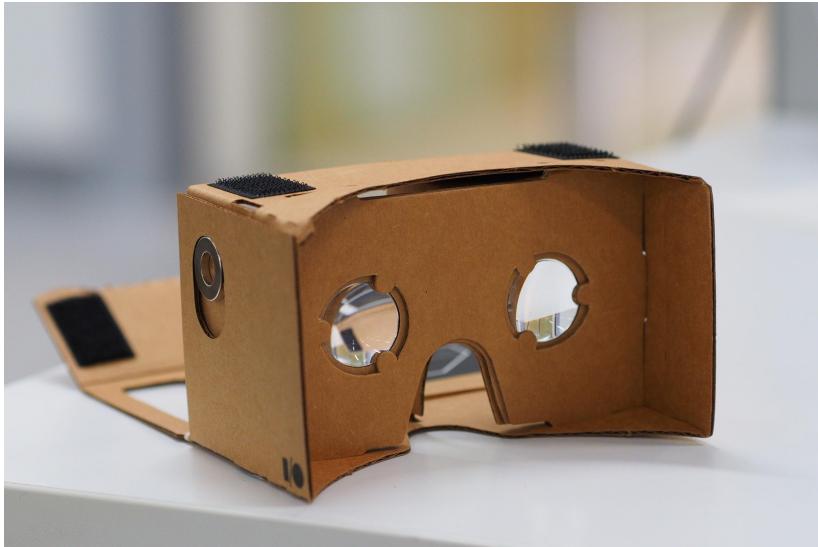
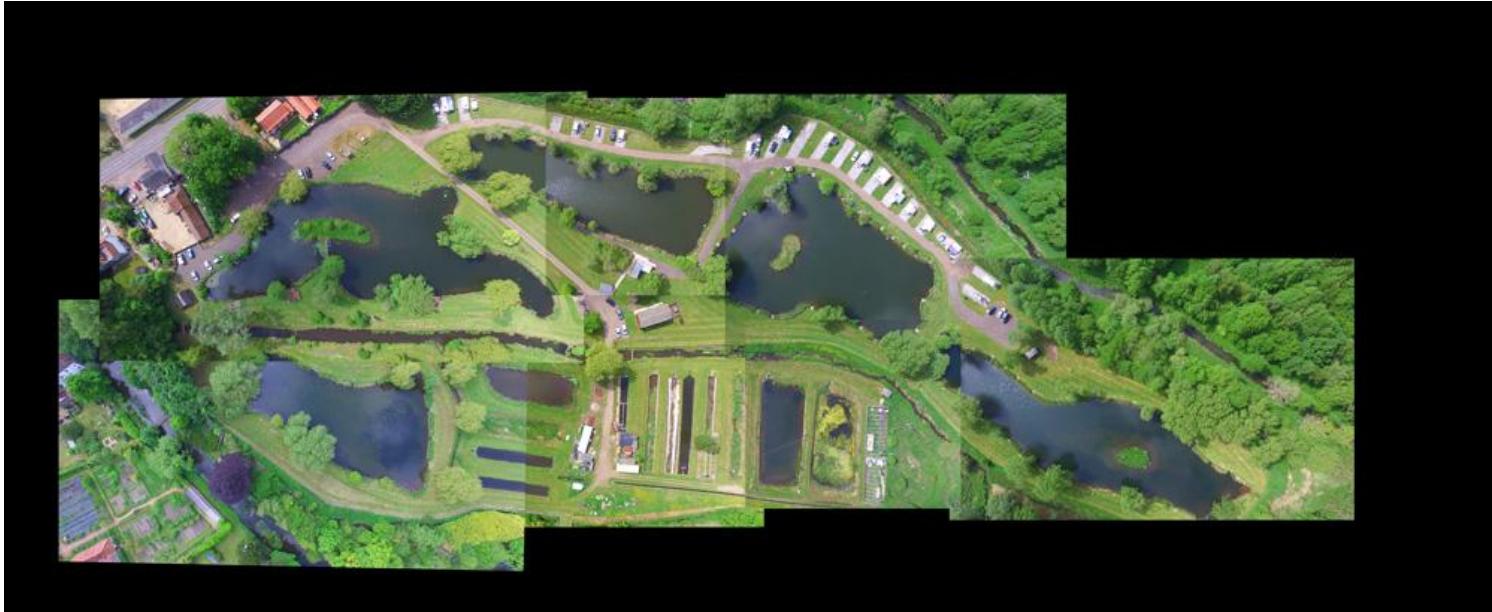


Image Alignment | Applications

Aerial mapping, Surveillance, Drones!



[Link](#)

Image Alignment | Applications

Image search, Search in image



[[OpenCV](#)]

Image Alignment | Applications

Video stabilization

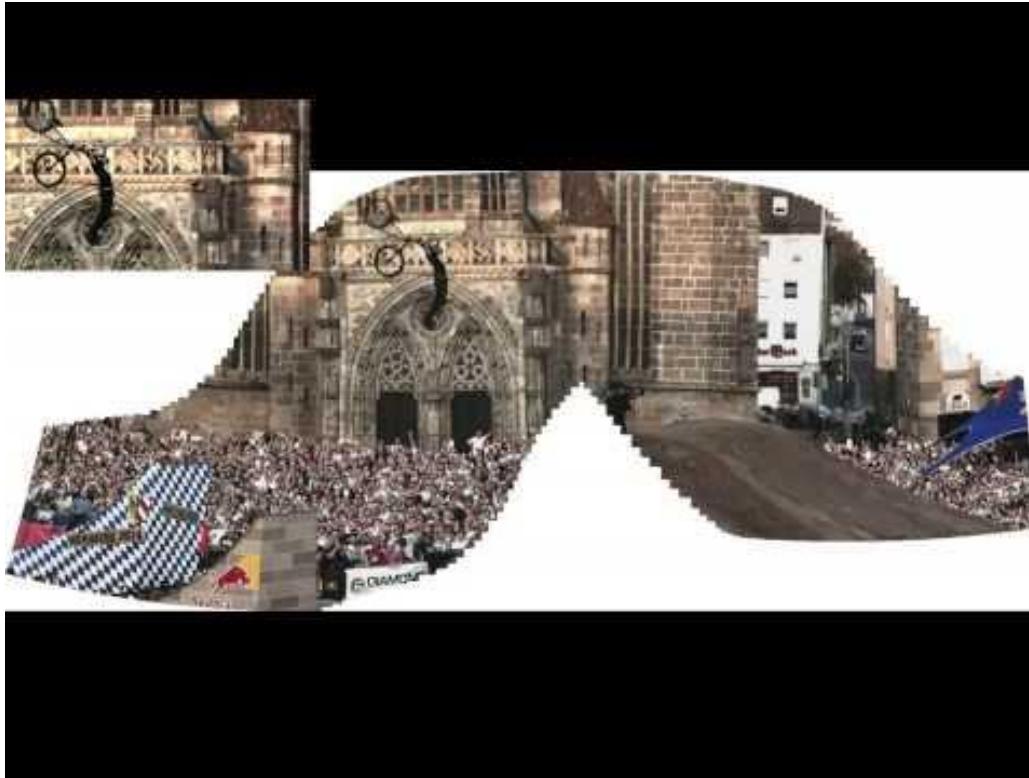


Image Alignment | Challenges

Is image alignment hard?

Between any two images we may have:

- Little overlap
- Occlusion
- Clutter (confusing features)
- Intensity changes
- Distortions
- Different hardware, optics
- ...

→ Alignment should be **robust** to a lot of variance.

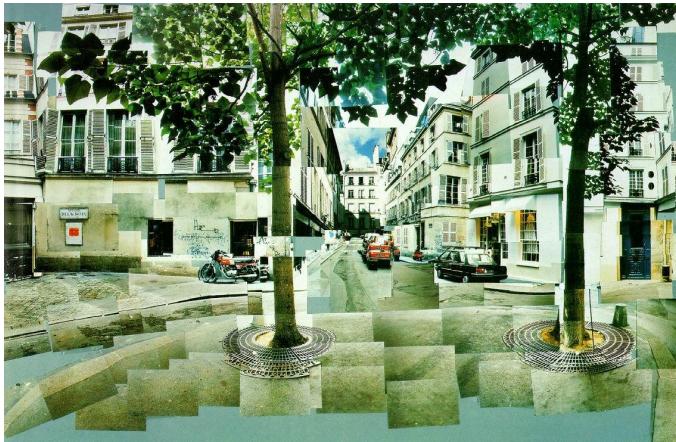


Image Alignment | Panography

What if we were aligning **real paper images**?

- We can **only rotate and translate**.

This was done: Panography.



David Hockney, 1986

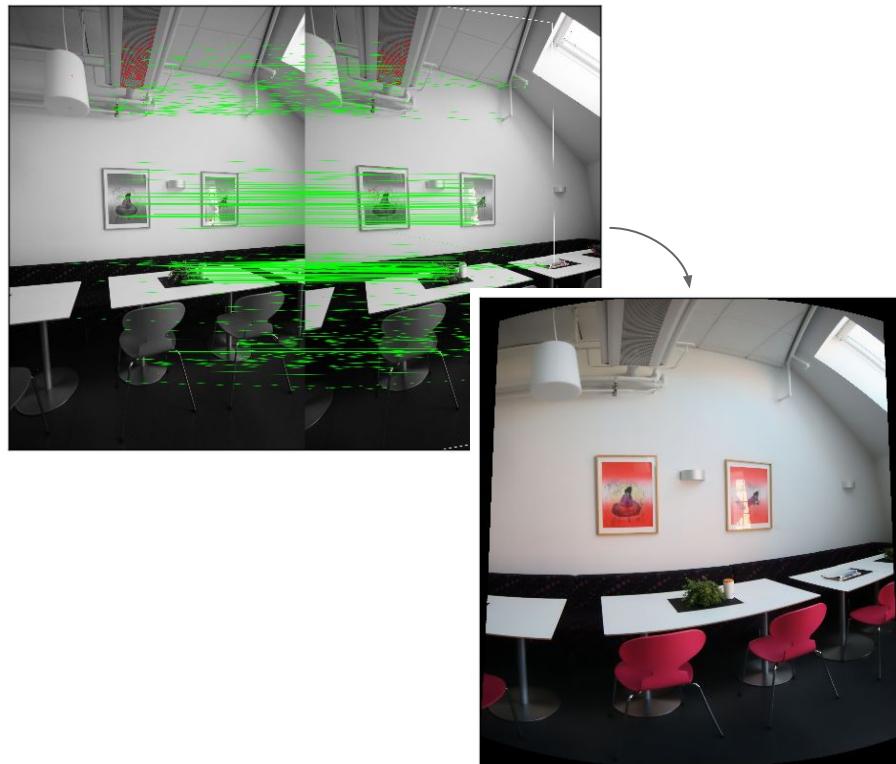
Feature Alignment

Instead of looking at pixels, we look at **interest points**.

- Less data to process
- More potential noise (from false matching)

Method overview:

1. Find key points and feature descriptors (Last time)
2. Match descriptors (we saw today)
3. Apply an overall geometric configuration (Today)
 - Affine, Homography
4. Find optimal parameters (Today)
 - Fitting, LSQ, robust LSQ
5. Blend



Alignment as a Fitting Problem

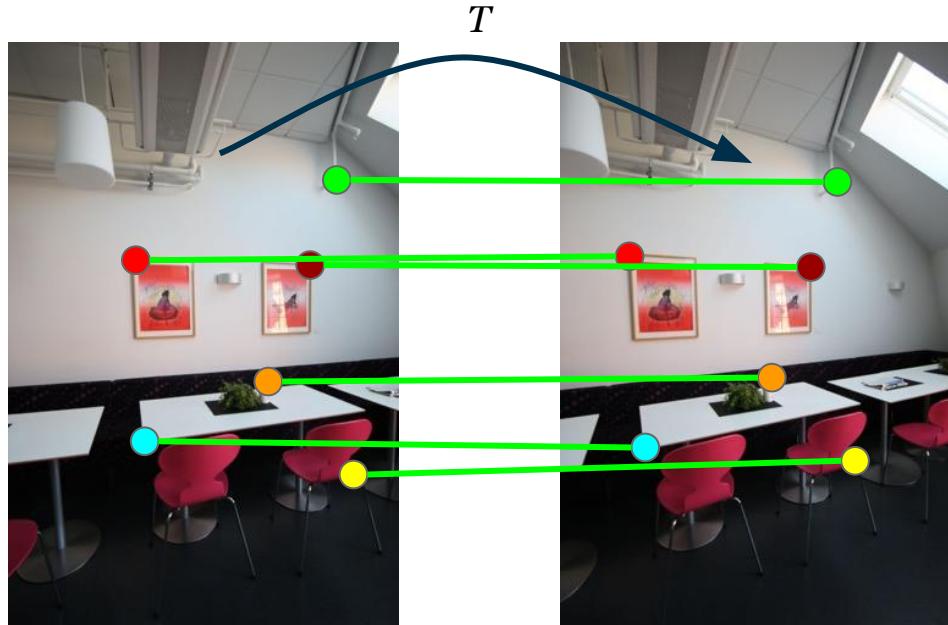
Over the feature matching:

- Design a **goodness of fit measure**

$$\hat{T} = \operatorname{argmin}_T f(T(X_1), X_2)$$

$$f(T(X_1), X_2) = \sum_i \|T(x_{1i}) - x_{2i}\|$$

- Design an **optimization scheme**
 - Avoid local minima
 - Robust to outliers



Simple Line Fitting: Least Squares

Input: Points $X = [(x_1, y_1), \dots, (x_n, y_n)]$

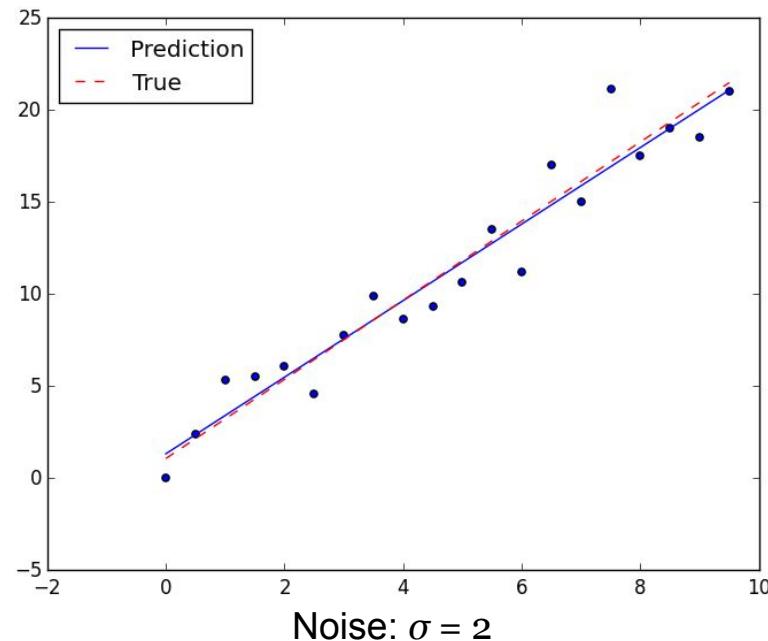
Output: (Best) Line through all points $y = \hat{a}x + \hat{b}$

$$a, b = \operatorname{argmin}_{a,b} E(X) = \sum_{i=1}^n (y_i - ax_i - b)^2$$

$$\begin{aligned} E(X) &= \sum_{i=0}^n \left\| (x_i, 1) \begin{bmatrix} a \\ b \end{bmatrix} - y_i \right\|^2 = \|Ax - y\|^2 \\ &= y^T y - 2(Ax)^T y + (Ax)^T (Ax) \end{aligned}$$

$$\frac{\partial E}{\partial x} = -2A^T y + 2A^T Ax = 0$$

$$A^T Ax = A^T y \Rightarrow \mathbf{x} = (A^T A)^{-1} A^T y$$



Least Squares

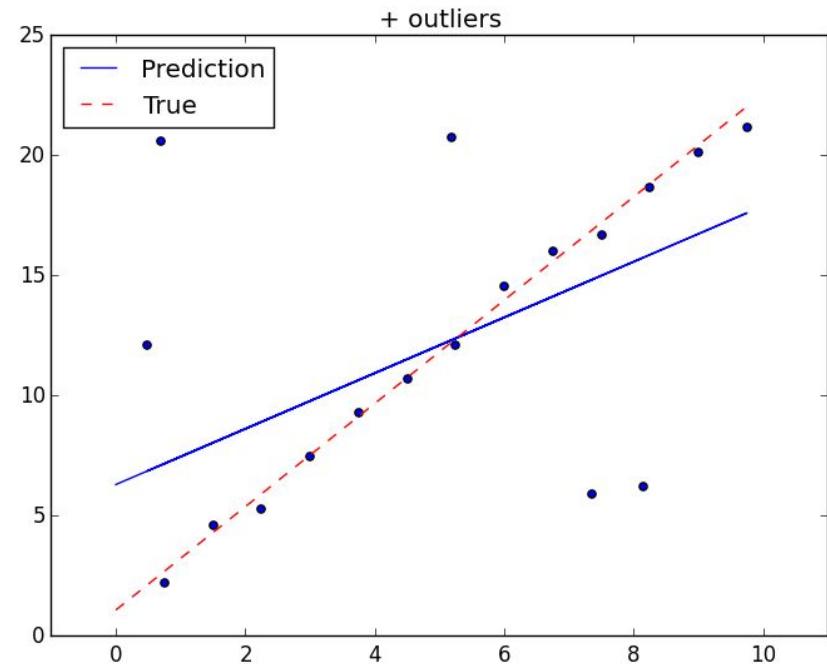
Good:

- Linear (fast, closed form)

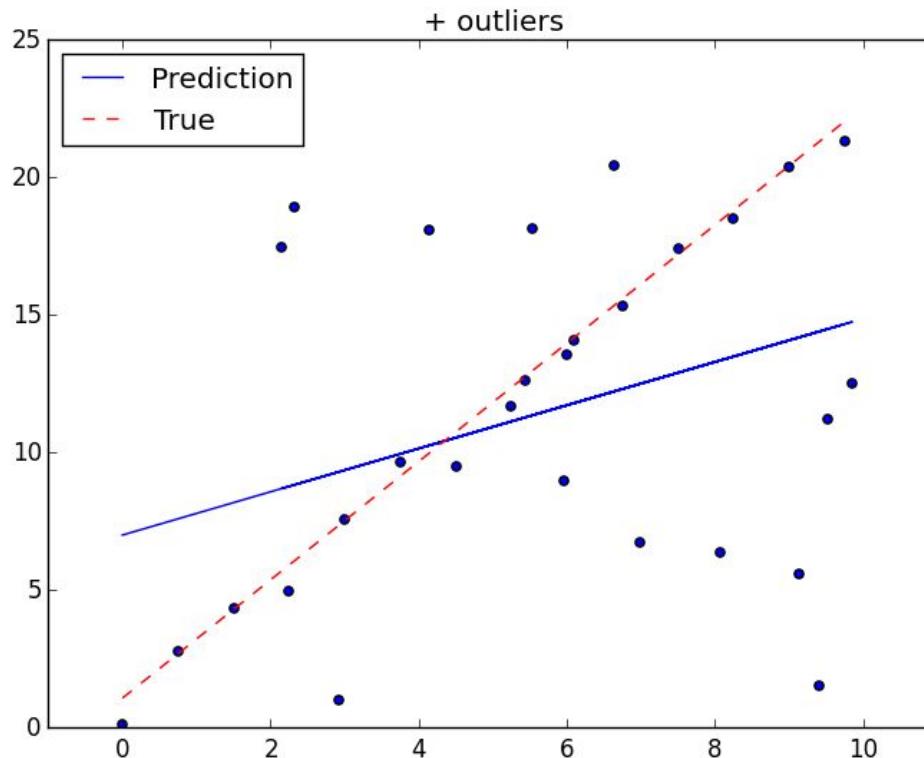
Bad:

- Sensitive to outliers (not robust)
- Only detects a single line

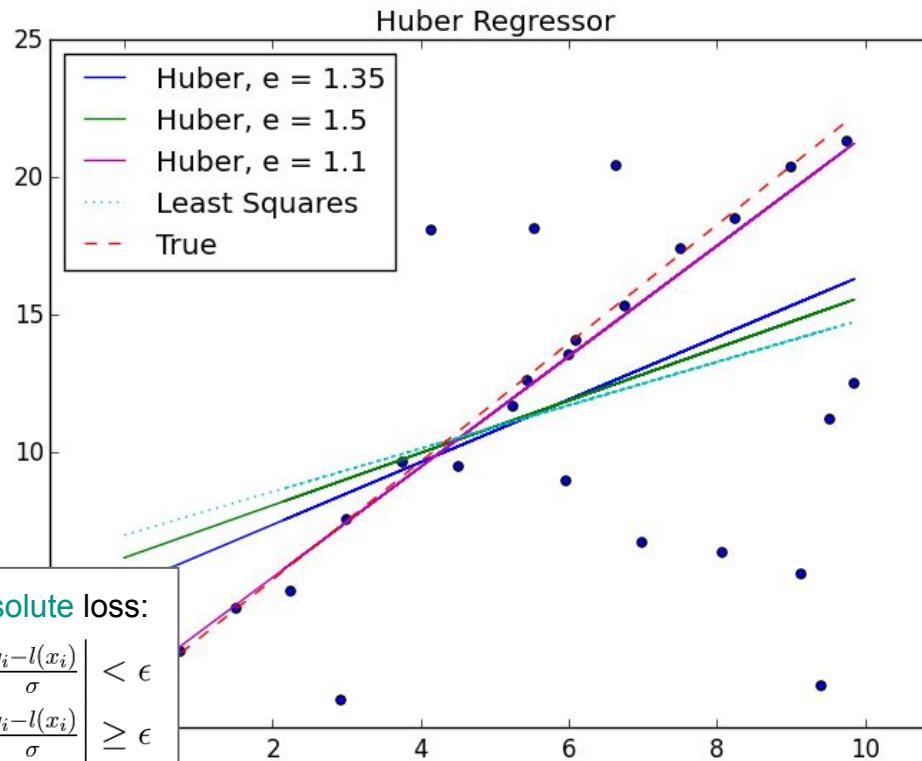
→ Use a **robust** estimator.



Robust Least Squares



Robust Least Squares



RANSAC

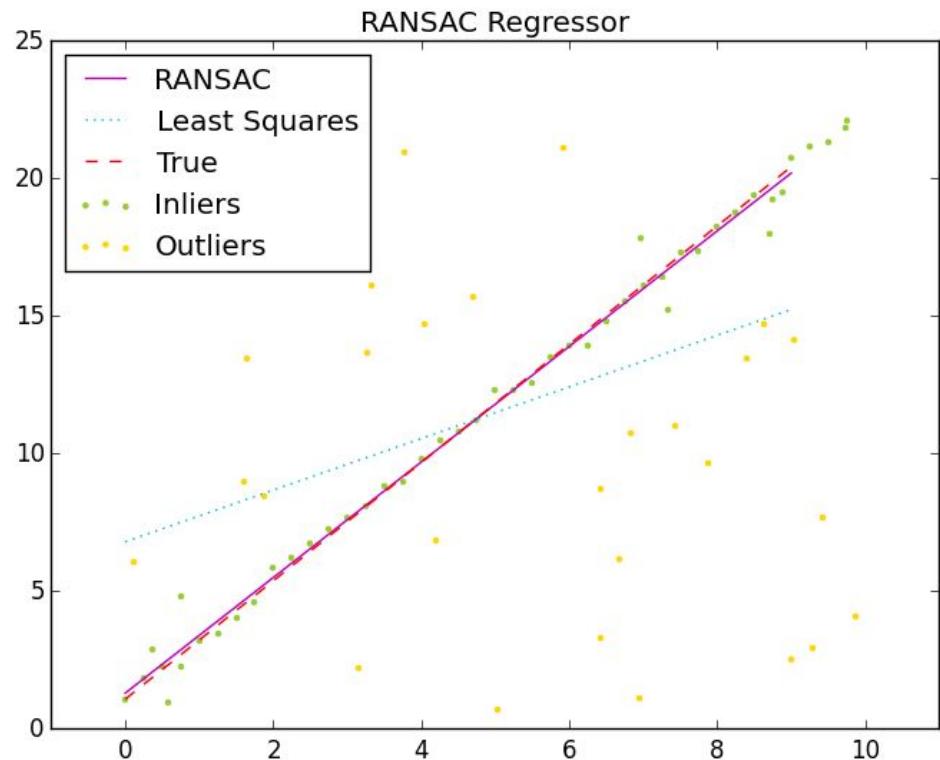
RAndom SAmples Consensus

So far - use *all* the points for matching & minimize the effect of outliers (robust estimators)

RANSAC idea: use only a *subset* of the points, discard outliers from the model fitting.

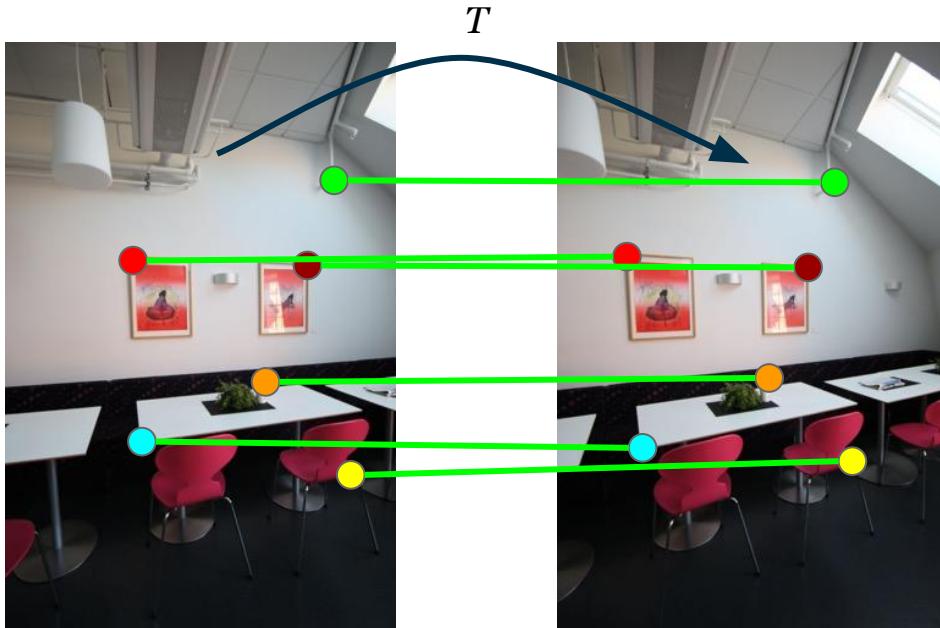
Method:

1. Choose a minimal set of samples
2. Estimate parameters (e.g. with LSq)
3. Count the number of inliers
4. Repeat 1-3 n times, pick best configuration
5. Refine the solution using only inliers



Questions?

Back to Alignment | Affine Transform



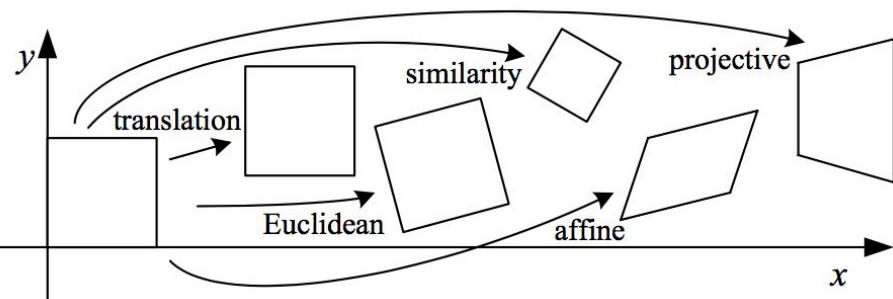
Let's assume T is **Affine**:

$$\begin{pmatrix} x'_i \\ y'_i \end{pmatrix} = \begin{pmatrix} m_1 & m_2 \\ m_3 & m_4 \end{pmatrix} \begin{pmatrix} x_i \\ y_i \end{pmatrix} + \begin{pmatrix} t_x \\ t_y \end{pmatrix}$$

Or in Homogeneous coordinates:

$$\begin{pmatrix} x'_i \\ y'_i \\ w'_i \end{pmatrix} = \begin{pmatrix} m_1 & m_2 & t_x \\ m_3 & m_4 & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_i \\ y_i \\ 1 \end{pmatrix}$$

2D Planar Transformations | Affine



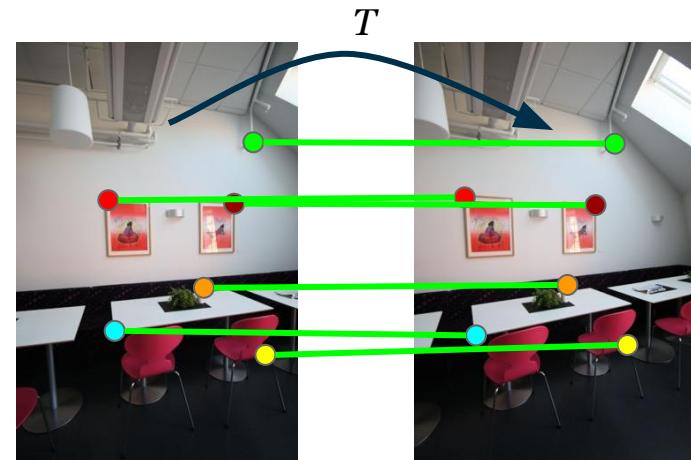
Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

2D **Affine** transformation:

- Parallel lines and planes remain parallel, but not angles or lengths.
- 6 DoF: 4 for rotation, scale and shear + 2 for translation
- Still linear → we can use Linear Least Squares to solve, and possibly add robustness to outliers

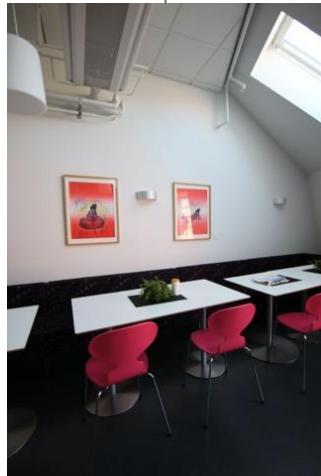
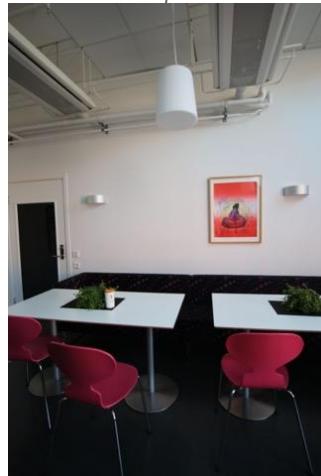
Affine Alignment

$$E(X) = \sum_i^n \left\| \begin{pmatrix} x'_i \\ y'_i \end{pmatrix} - \begin{pmatrix} m_1 & m_2 \\ m_3 & m_4 \end{pmatrix} \begin{pmatrix} x_i \\ y_i \end{pmatrix} - \begin{pmatrix} t_x \\ t_y \end{pmatrix} \right\|_{L_2}$$
$$= \left\| \begin{pmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 1 \\ 0 & 0 & x_i & y_i & 1 & 1 \\ \dots & & & & & \end{pmatrix} \begin{pmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_x \\ t_y \end{pmatrix} - \begin{pmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{pmatrix} \right\|_{L_2}$$



A non-homogeneous system of linear equations, with 6 DoF.

Affine Alignment | Result



Somewhat aligned,
But it will simply not do.

Homographies

What can we do besides Affine? a **Homography**.

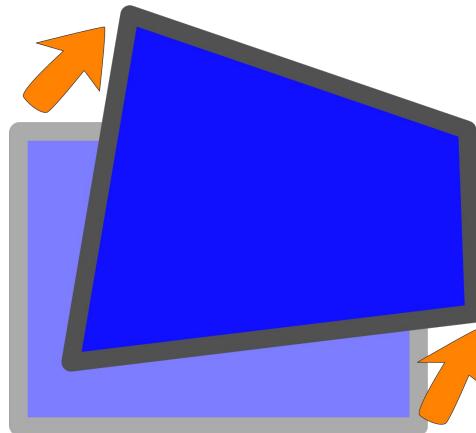
- It does not preserve parallel lines.
- Still linear - we can use the same exact fitting process.

$$H = \begin{pmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & 1 \end{pmatrix}$$

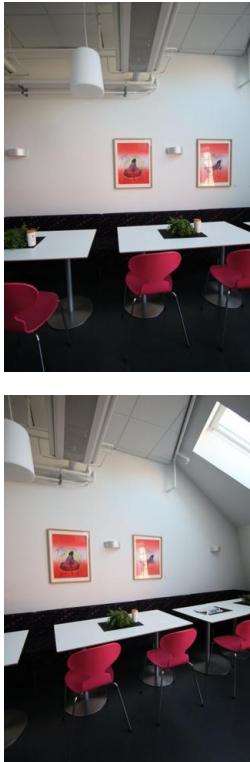
We added two more elements - what do they do?

Allow us to model **perspective** transformation.

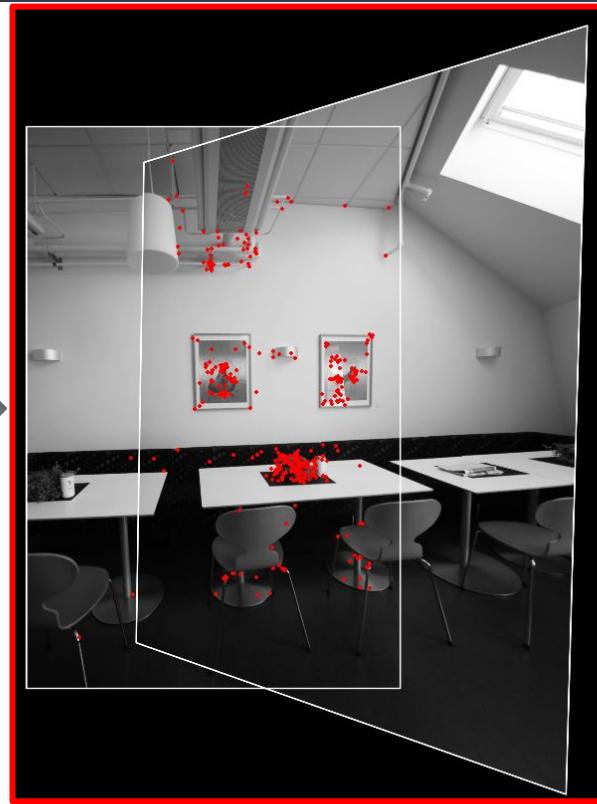
$$\begin{pmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \Rightarrow \begin{pmatrix} \frac{h_1x+h_2y+h_3}{h_7x+h_8y+1} \\ \frac{h_4x+h_5y+h_6}{h_7x+h_8y+1} \\ 1 \end{pmatrix}$$



Homography-based Alignment



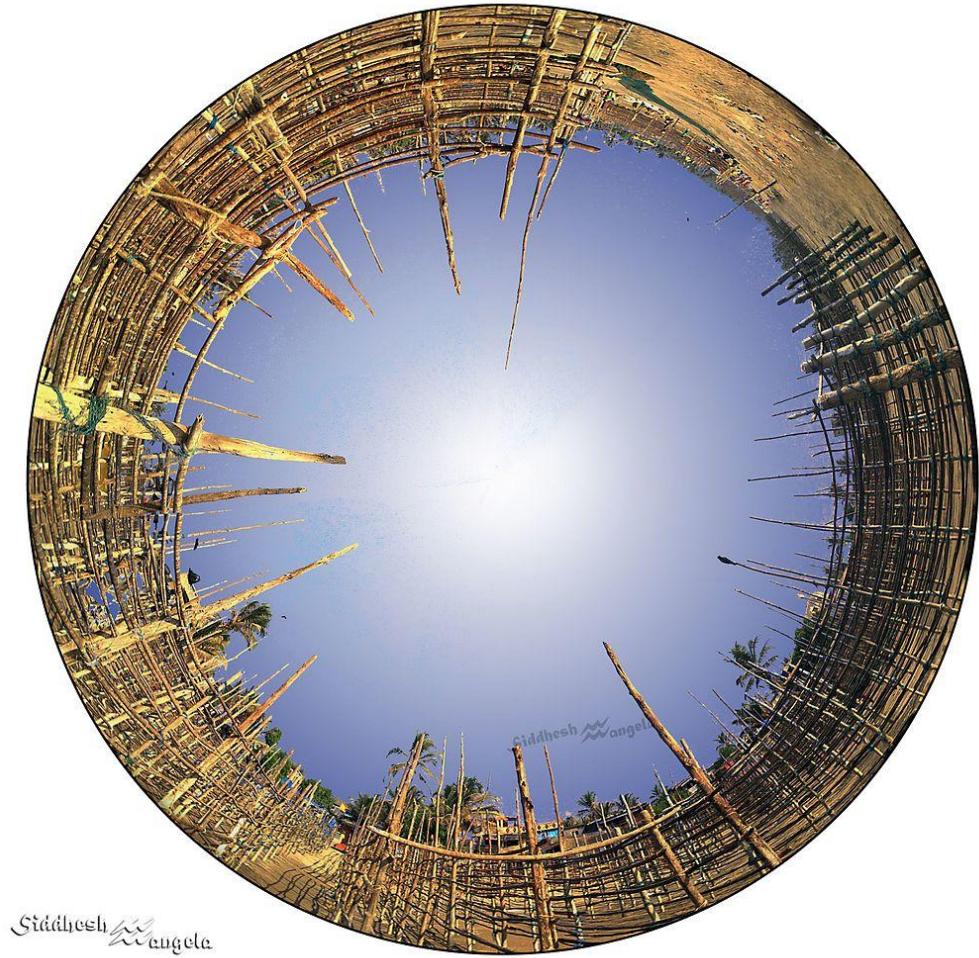
Affine? Meh.



Homography!
Yay!

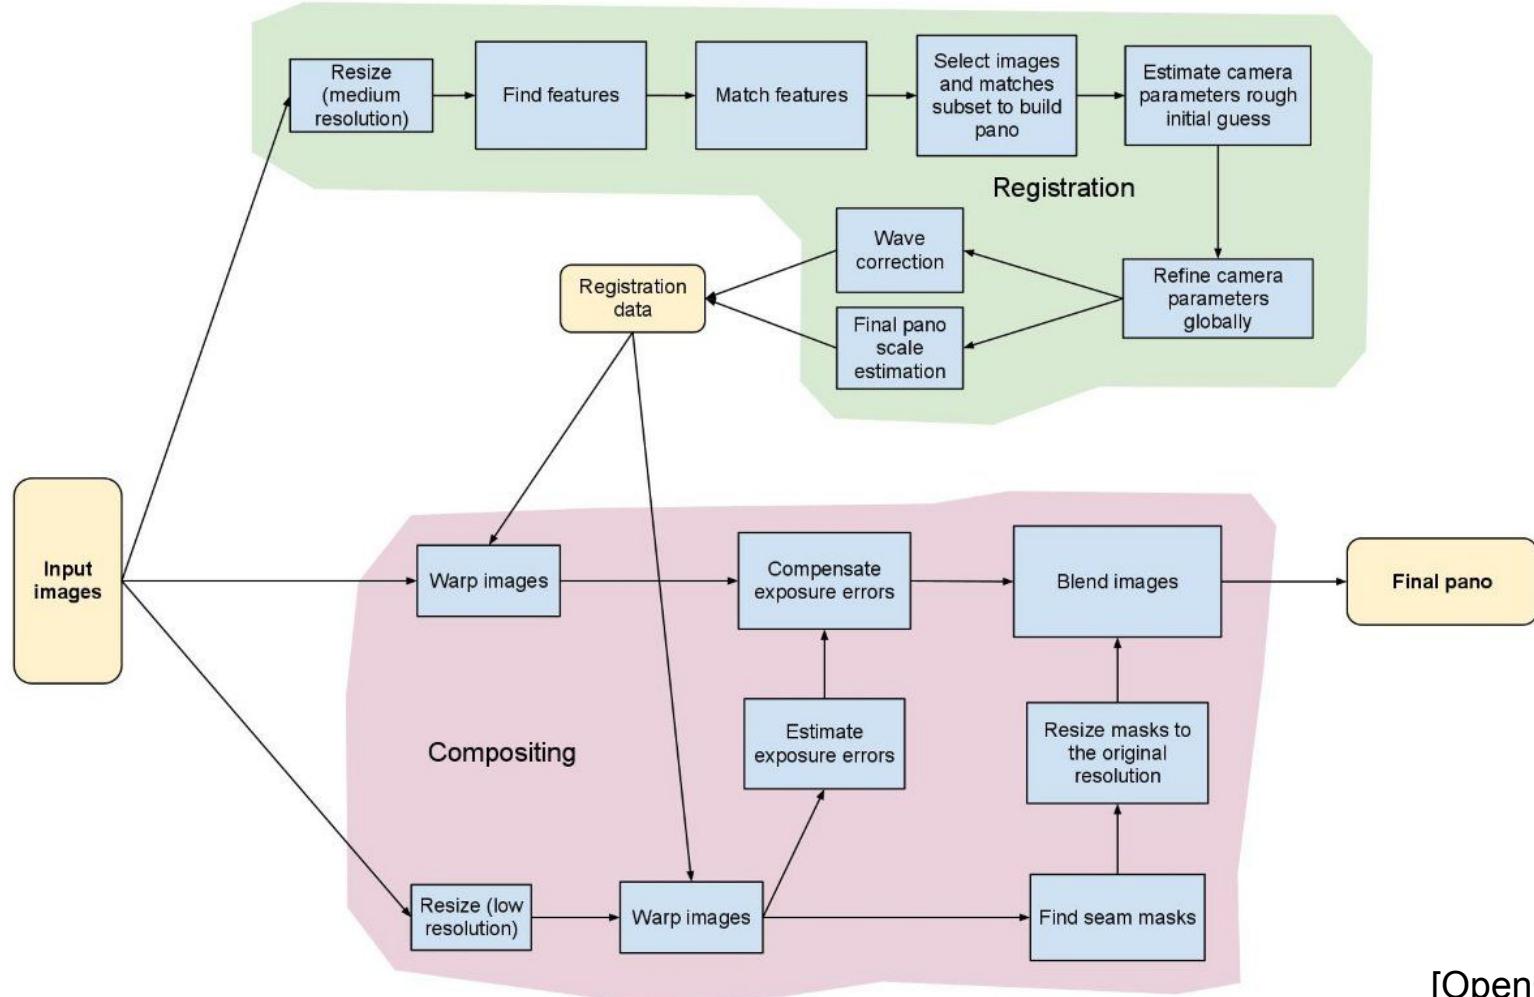
Questions?

Panoramas



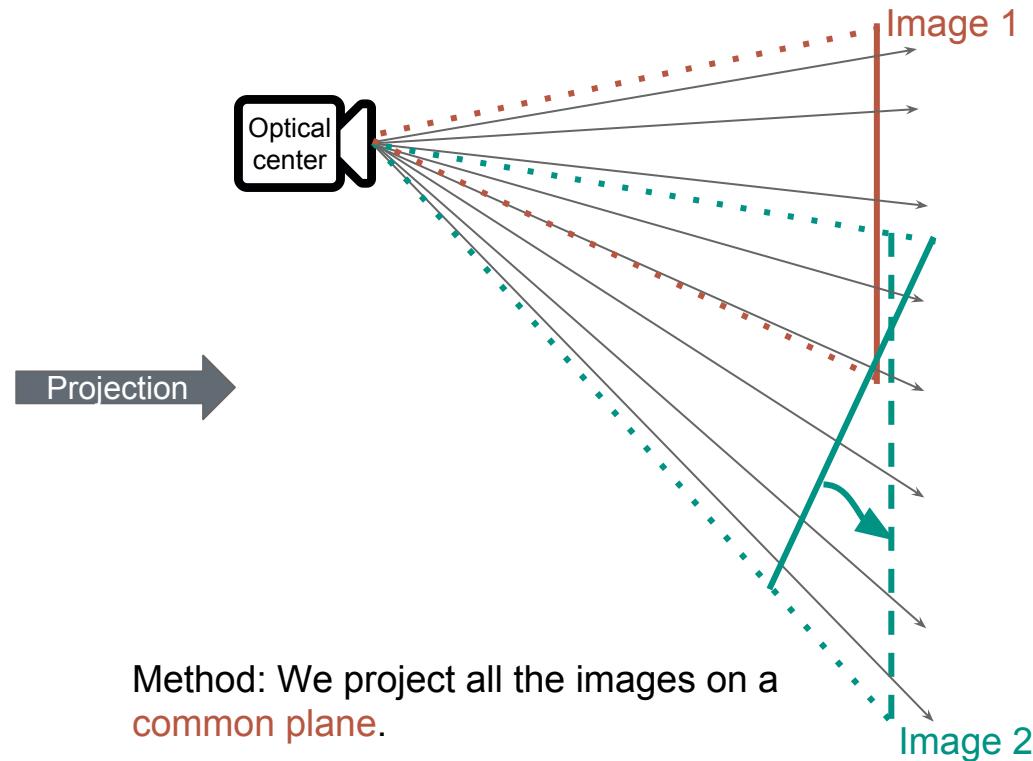
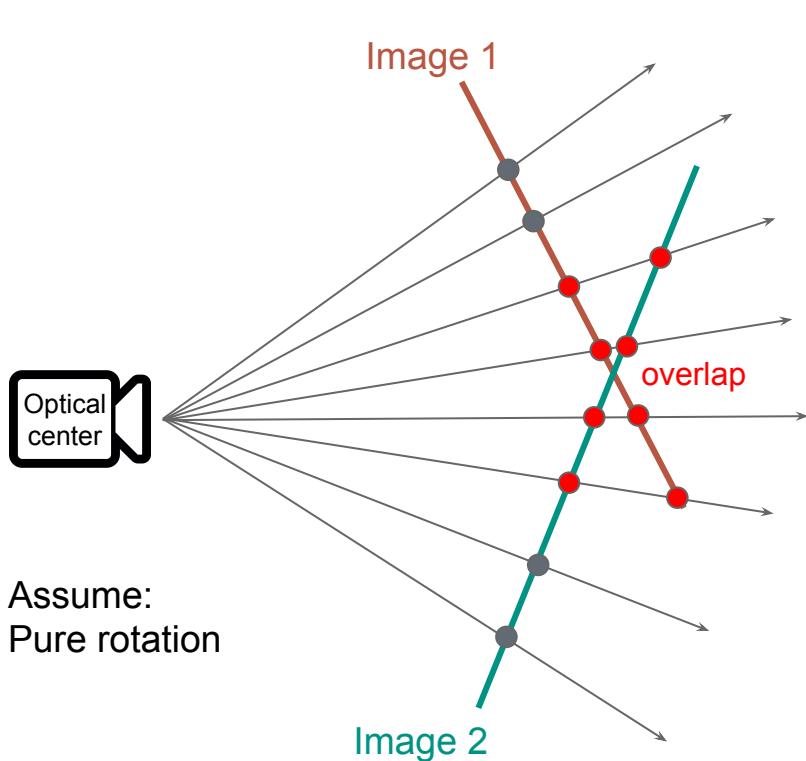
Ciddhesh
angle

Ciddhesh
angle

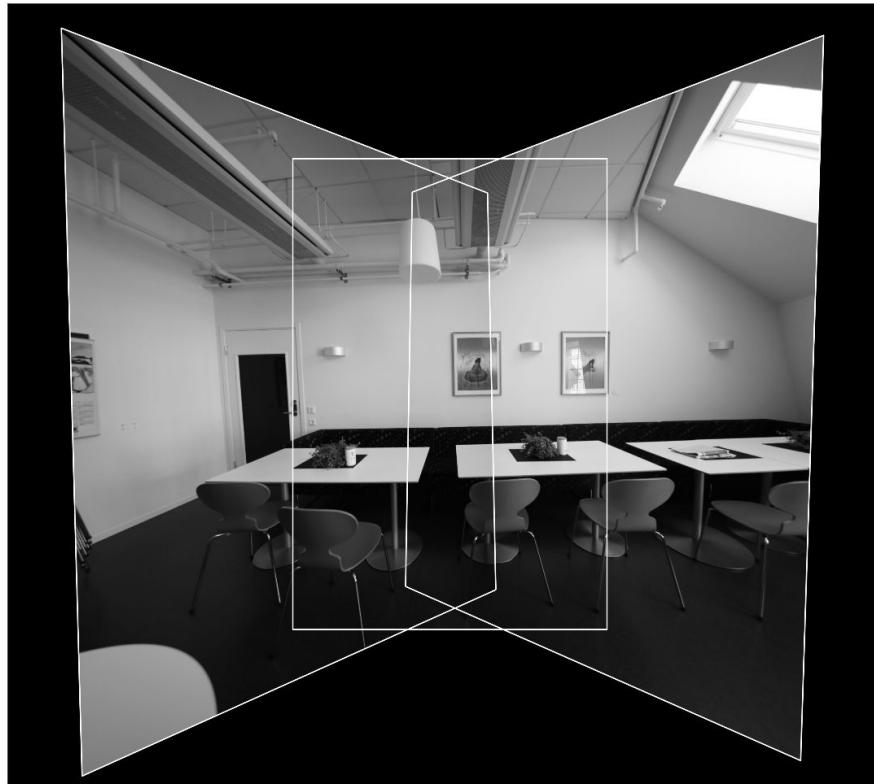


[OpenCV]

Mosaics: Geometry

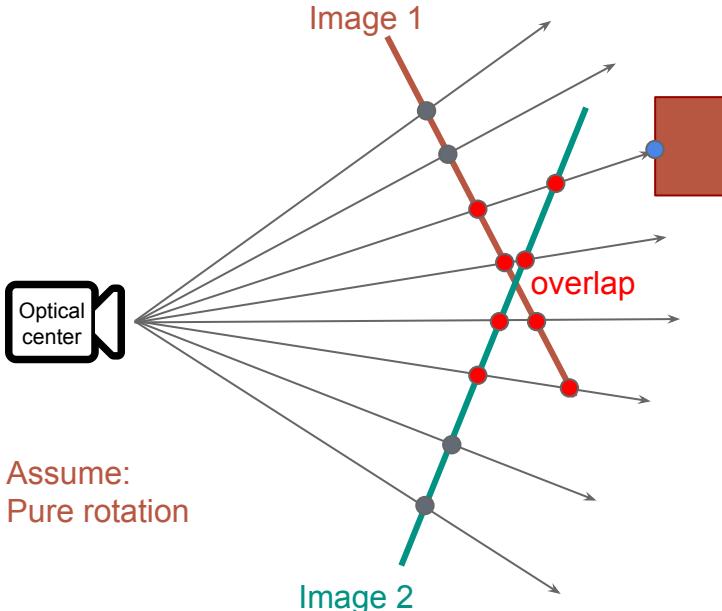


Projective Panorama



Panorama Geometry

What happens to the overlapping points?



$$\text{In camera 2: } \begin{pmatrix} X_2 \\ Y_2 \\ Z_2 \end{pmatrix} = K_2^{-1} \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix}$$

$$\text{Move to camera 1: } \begin{pmatrix} X_1 \\ Y_1 \\ Z_1 \end{pmatrix} = \underbrace{R_1 R_2^T K_2^{-1}}_{\text{Undo 2, Apply 1}} \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix}$$

$$\text{Project on image 1: } \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix} \approx \underbrace{K_1 R_1 R_2^T K_2^{-1}}_{3 \times 3 \text{ Homography}} \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix}$$

⇒ Homography is:

- (1) Undo projection 2,
- (2) rotate from cam 2 to cam 1,
- (3) apply projection 1

A 360° Panorama With Homographies?



eek.

Not really.

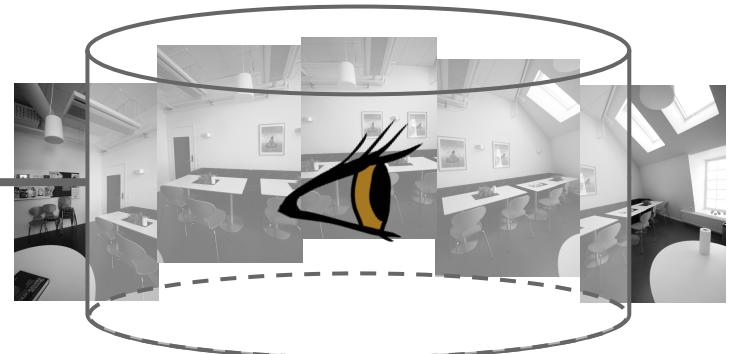
360° Panoramas



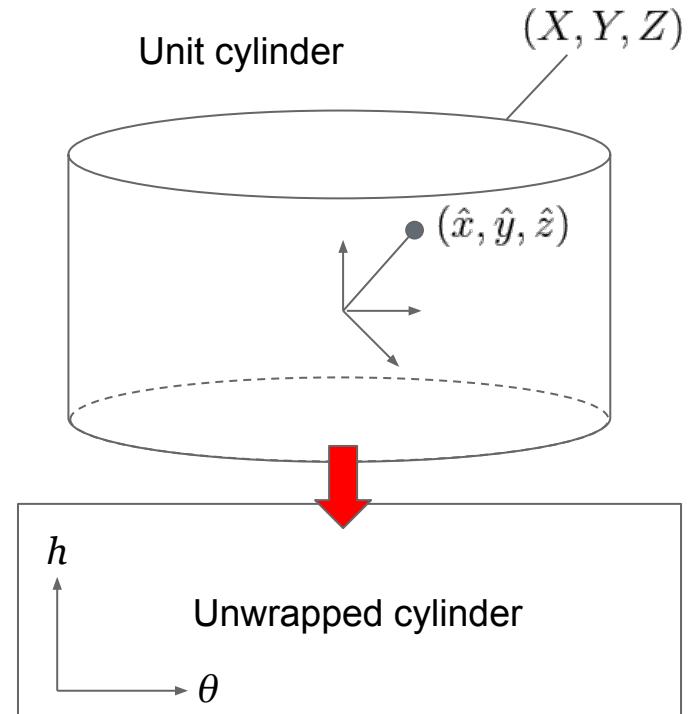
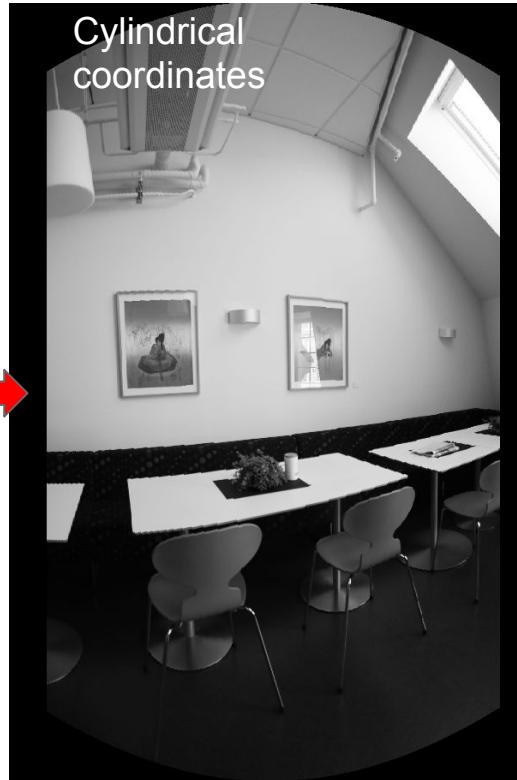
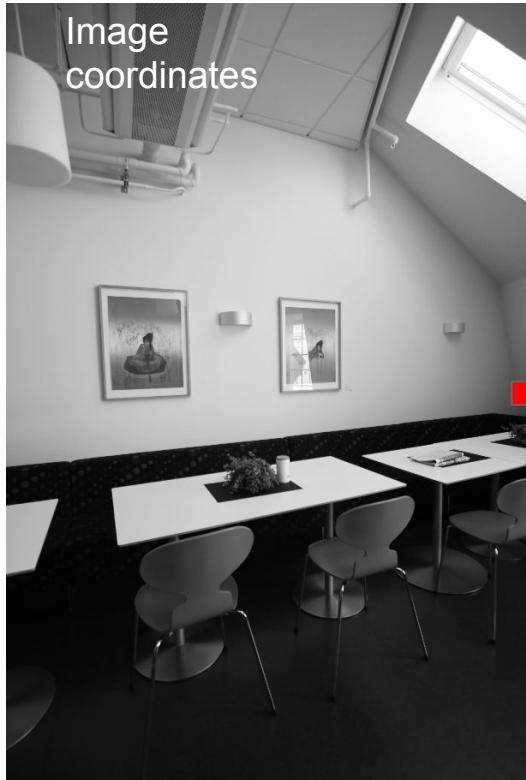
View from inside the cylinder.



Imagine the images are on the wall of a cylinder...



A 360° Panorama | Cylindrical Warping



A 360° Panorama | Cylindrical Warping

Easier to go the inverse, from cylinder to image:

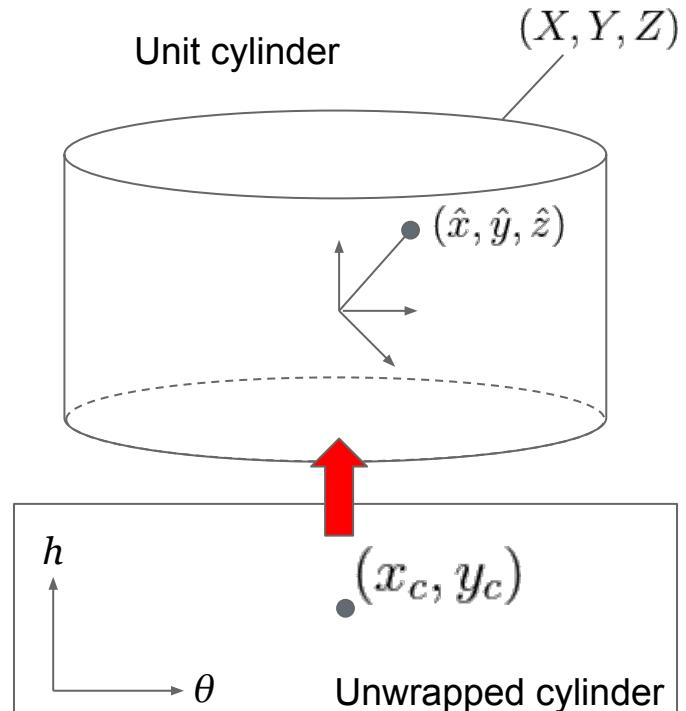
$$\theta = (x_{cyl} - x_c)/f$$

$$h = (y_{cyl} - y_c)/f$$

$$\hat{x} = \sin \theta$$

$$\hat{y} = h$$

$$\hat{z} = \cos \theta$$



Aligning in Cylindrical Coordinates

We **rotated** the camera by θ .

How does this affect the cylindrical image?



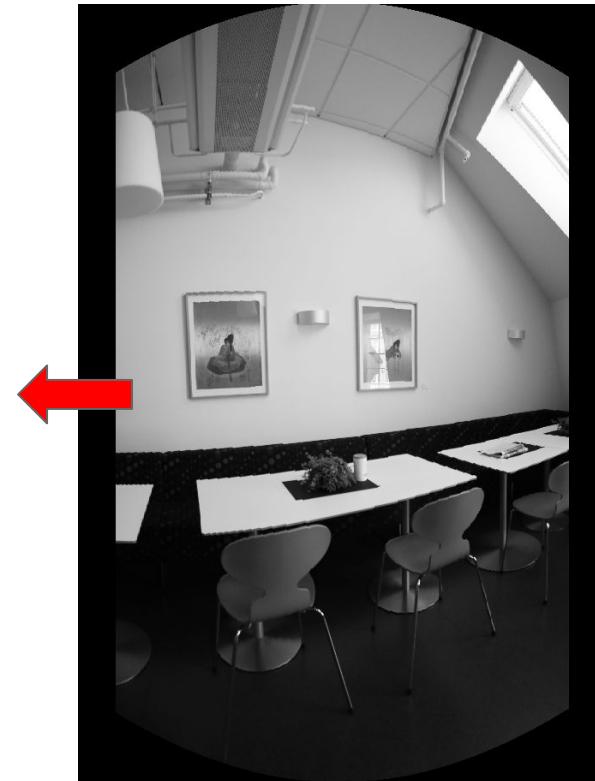
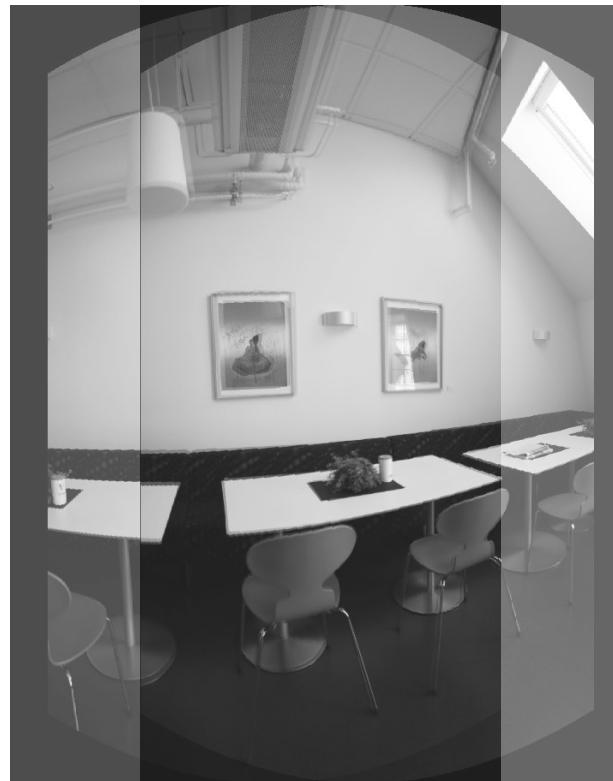
Aligning in Cylindrical Coordinates

We **rotated** the camera by θ .

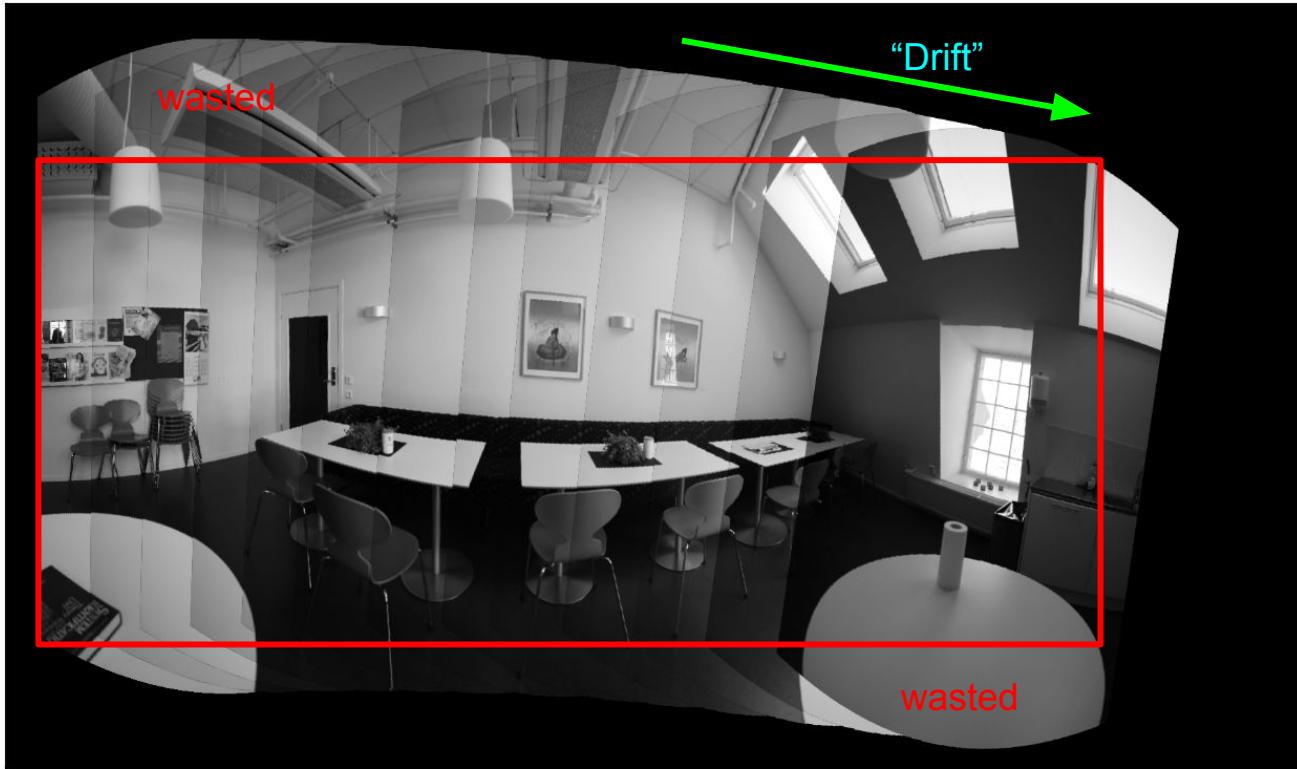
How does this affect the cylindrical image?

It's a (simple) **translation**.

Aligning cylindrical images
needs only a translational model
of 2 DoF.



How to Cut?



Questions?

HW2: Image Stitching

Your goal is to create 2 panoramas:

1. Using homographies and perspective warping on a common plane (3 images).
2. Using cylindrical warping (many images).

In both options you should:

- Extract features and descriptors
- Match features
- Calculate the best model parameters
- Bonus (!!): Use your Laplacian Blending code to stitch the images together nicely

You will be given sample code to

- convert an image to cylindrical coordinates
- calculate affine or homography transform
- skeleton panorama maker

The rest is up to you...